

PROJECT EVALUATION RUBRIC

Scoring Guide for Technical Project Assessment

OVERVIEW

This rubric is used to evaluate technical project implementations, specifically for backend systems involving AI/LLM integration, RAG implementations, and case study submissions. Each parameter is scored on a scale of 1-5. The final score is a weighted average of all parameters.

SCORING PARAMETERS

1. CORRECTNESS (PROMPT & CHAINING) (Weight: 30%)

Description:

Evaluate whether the project correctly implements the required functionality, including prompt design, LLM chaining logic, RAG context injection, and overall system correctness.

Scoring Scale:

1 - NOT IMPLEMENTED

- Core functionality is missing
- No LLM integration or it doesn't work
- No RAG implementation
- Requirements not met
- System doesn't run or crashes immediately

2 - MINIMAL ATTEMPT

- Basic structure present but mostly incomplete
- LLM integration exists but doesn't work properly
- No actual chaining or very broken chaining
- RAG attempted but not functional
- Many core requirements missing
- Significant bugs prevent proper operation

3 - WORKS PARTIALLY

- Core functionality partially implemented
- LLM integration works for simple cases
- Basic chaining logic present but limited
- RAG retrieves context but integration is weak
- Some requirements met, others missing
- Works for happy path but fails on edge cases

4 - WORKS CORRECTLY

- All core functionality implemented and working
- Proper LLM integration with good prompts
- Chaining logic works reliably
- RAG effectively retrieves and injects context
- All major requirements met
- Handles most edge cases
- Minor issues or missing optimizations

5 - FULLY CORRECT + THOUGHTFUL

- Excellent implementation of all requirements
- Sophisticated prompt design with clear structure
- Robust multi-stage LLM chaining
- Effective RAG with proper chunking and retrieval
- All requirements exceeded
- Thoughtful design decisions
- Handles edge cases elegantly
- Goes beyond basic requirements with intelligent solutions

2. CODE QUALITY & STRUCTURE (Weight: 25%)

Description:

Evaluate the cleanliness, modularity, reusability, and testability of the code.

Scoring Scale:

1 - POOR

- Messy, unorganized code
- No structure or separation of concerns
- Hard-coded values everywhere
- No comments or documentation
- Copy-pasted code with lots of duplication
- Very difficult to understand or modify

2 - SOME STRUCTURE

- Basic organization present
- Some separation of concerns
- Limited modularity
- Minimal comments
- Some code duplication
- Difficult to extend or maintain

3 - DECENT MODULARITY

- Good basic structure
- Clear separation of concerns
- Functions/classes are reasonably sized
- Some reusable components
- Basic error handling
- Code is readable with some comments
- Could use more abstraction
- No tests or very minimal testing

4 - GOOD STRUCTURE + SOME TESTS

- Clean, well-organized code
- Strong separation of concerns
- Good modularity and abstraction
- Reusable components
- Clear naming conventions
- Well-commented where needed
- Configuration separated from code
- Some unit tests or integration tests
- Easy to understand and extend

5 - EXCELLENT QUALITY + STRONG TESTS

- Exceptional code quality
- Excellent architecture and design patterns
- Highly modular and reusable
- Comprehensive error handling
- Clear, self-documenting code
- Configuration-driven
- Comprehensive test coverage (unit + integration)
- Follows industry best practices
- Production-ready code quality

3. RESILIENCE & ERROR HANDLING (Weight: 20%)

Description:

Evaluate how well the system handles failures, implements retry logic, manages long-running processes, and controls randomness/nondeterminism.

Scoring Scale:

1 - MISSING

- No error handling
- System crashes on errors

- No retry logic
- No timeout protection
- Doesn't handle API failures
- LLM randomness not controlled

2 - MINIMAL

- Basic try-catch blocks
- Generic error messages
- No retry mechanisms
- Limited error logging
- Fails ungracefully on API errors
- No consideration for LLM stability

3 - PARTIAL HANDLING

- Some error handling present
- Basic retry logic for some operations
- Some timeout protection
- Logs errors but limited detail
- Handles some failure scenarios
- Basic LLM temperature control
- Inconsistent error handling across the system

4 - SOLID HANDLING

- Good error handling throughout
- Retry logic with exponential backoff
- Proper timeout protection
- Detailed error logging
- Handles most failure scenarios gracefully
- LLM temperature controlled for consistency
- Graceful degradation on failures
- Clear error messages to users

5 - ROBUST, PRODUCTION-READY

- Comprehensive error handling
- Intelligent retry strategies with backoff
- Circuit breaker pattern for external services
- Detailed logging and monitoring
- Handles all foreseeable failure scenarios
- Multiple strategies for LLM consistency (temperature + validation)
- Graceful degradation with partial results
- Dead letter queue for failed jobs
- Health checks and self-healing capabilities
- Production-grade resilience patterns

4. DOCUMENTATION & EXPLANATION (Weight: 15%)

Description:

Evaluate the quality of README, setup instructions, code comments, and explanations of design decisions and trade-offs.

Scoring Scale:

1 - MISSING

- No README or very minimal
- No setup instructions
- No explanation of how to run the project
- No documentation of design decisions
- No comments in code

2 - MINIMAL

- Basic README present
- Incomplete setup instructions
- Missing many details
- No explanation of design choices
- Minimal or no code comments
- Difficult to get started

3 - ADEQUATE

- README covers basics
- Setup instructions present but may lack detail
- Basic explanation of project structure
- Some design decisions mentioned
- Basic code comments
- Possible to get started with some effort
- Missing some important details

4 - CLEAR

- Comprehensive README
- Clear, step-by-step setup instructions
- Good explanation of architecture
- Design decisions documented
- Code is well-commented
- Easy to get started and understand
- Includes examples of usage
- Mentions some trade-offs

5 - EXCELLENT + INSIGHTFUL

- Outstanding documentation
- Multiple documentation files (README, SETUP, ARCHITECTURE)
- Crystal-clear setup instructions with troubleshooting
- Detailed explanation of all design decisions
- Thorough discussion of trade-offs and alternatives considered
- Excellent code comments where needed
- API documentation (e.g., Swagger/OpenAPI)
- Diagrams or visualizations of architecture
- Examples and usage scenarios
- Rationale for technology choices
- Known limitations and future improvements discussed

5. CREATIVITY / BONUS (Weight: 10%)

Description:

Evaluate additional features, innovative solutions, and improvements beyond the basic requirements.

Scoring Scale:

1 - NONE

- Strictly minimum requirements only
- No additional features
- No innovation or creative solutions
- Basic implementation throughout

2 - VERY BASIC

- One or two minor extras
- Basic implementation of bonuses
- Limited additional value
- Extras don't add much to the project

3 - USEFUL EXTRAS

- Several additional features that add value
- Examples:
 - Additional API endpoints (health check, statistics)
 - Basic monitoring or logging
 - Docker setup for easy deployment
 - Environment-based configuration
- Thoughtful but not exceptional additions

4 - STRONG ENHANCEMENTS

- Multiple valuable additions beyond requirements

- Examples:
 - Comprehensive API with extra endpoints
 - Advanced monitoring (e.g., Flower for Celery)
 - Full Docker Compose setup
 - Multiple LLM provider support
 - Caching layer
 - API rate limiting
 - Admin dashboard
- Demonstrates strong engineering thinking
- Adds significant value to the project

5 - OUTSTANDING CREATIVITY

- Exceptional additional features and innovations
- Examples:
 - Multi-model ensemble approach
 - Real-time streaming updates (WebSocket/SSE)
 - Advanced caching with invalidation
 - A/B testing framework for prompts
 - Automated prompt optimization
 - Comprehensive monitoring and alerting
 - Authentication and authorization system
 - Deployment to cloud with CI/CD
 - Performance optimizations
 - Unique innovative solutions to problems
- Production-ready additional features
- Shows exceptional technical depth and breadth

CALCULATION METHOD

Final Project Score = Weighted Average of All Parameters

Example Calculation:

- Correctness: 4 (weight 0.30) = 1.20
- Code Quality: 4 (weight 0.25) = 1.00
- Resilience: 4 (weight 0.20) = 0.80
- Documentation: 5 (weight 0.15) = 0.75
- Creativity: 3 (weight 0.10) = 0.30

Final Score = $1.20 + 1.00 + 0.80 + 0.75 + 0.30 = 4.05$

INTERPRETATION GUIDE

Project Score Interpretation:

- 4.5 - 5.0: Exceptional work, exceeds expectations
- 4.0 - 4.4: Strong work, meets all expectations
- 3.5 - 3.9: Good work, meets most expectations
- 3.0 - 3.4: Adequate work, meets basic expectations
- 2.5 - 2.9: Below expectations, significant gaps
- Below 2.5: Well below expectations, major issues

EVALUATION GUIDELINES

1. Focus on implementation quality, not just functionality
2. Consider both correctness and code craftsmanship
3. Value resilience and production-readiness
4. Reward clear communication through documentation
5. Appreciate creative solutions and extras
6. Be fair but maintain high standards
7. Provide specific, actionable feedback

COMMON PITFALLS TO WATCH FOR

- Hard-coded credentials or API keys in code
- No error handling for external API calls
- Synchronous operations that should be asynchronous
- Missing or inadequate documentation
- Copy-pasted code without understanding
- No separation between configuration and code
- Overly complex solutions to simple problems
- No consideration for edge cases
- Missing tests entirely
- Poor naming conventions

EXCELLENCE INDICATORS

- Clean architecture with clear separation of concerns
- Comprehensive error handling with graceful degradation
- Well-designed prompts with context injection
- Effective use of RAG for improved accuracy
- Proper async processing for long-running tasks
- Extensive documentation with diagrams
- Test coverage for critical paths
- Production-ready deployment setup

- Thoughtful trade-off discussions
- Creative solutions that add value

NOTES FOR EVALUATORS

- Provide constructive feedback for improvement
- Highlight specific examples of good and bad practices
- Consider the candidate's experience level
- Focus on potential and learning ability
- Value clean, simple solutions over complex ones
- Encourage best practices and production thinking
- Recognize effort in documentation and testing
- Appreciate thoughtful design decisions